

CLAUDE

CLAUDE.md Working context for any AI assistant (Claude, Cursor, Copilot, others) operating on this repo. What this project is ClauseGuard is an evidence-grounded legislative risk analysis system for Ghana. It reads bills, splits them into clauses, classifies risks across eleven categories, and checks each clause against the 1992 Constitution of Ghana. Every output is backed by a cited excerpt. The primary user is a mid-level analyst at a Ghanaian industry association who has seven to fourteen days to produce a position paper on a newly gazetted bill. Secondary users are parliamentary committee staff, civil society organisations, and journalists. What this project is not It is not a legal opinion engine. It is not partisan. It does not predict parliamentary outcomes. It does not give advice. It surfaces evidence. The human decides. Owner Owura Bediako. Founder of Govaxis AI. Based in London, originally from Ghana. Writing rules These apply to every prose output the assistant produces for this project. README, reports, blog posts, LinkedIn posts, commit messages, code comments, docstrings. No em dashes. No semicolons. No colons in prose. Acceptable in code, tables, file paths, and time stamps. No markdown bold or italics in the body of reports unless explicitly requested. Headings are fine. Active voice by default. Passive only when it reads better. Short sentences. Two clauses maximum.

No AI filler vocabulary. No "leverage", "delve", "in today's rapidly evolving landscape", "it is important to note", "navigate the complexities". No three-item rhetorical lists. No negative-contrast constructions. "Not just X, but Y" is banned. State what the thing is, directly. No metaphors and no clichés. No unnecessary adjectives and adverbs. No output warnings or hedging notes at the end of responses. These rules apply across all the founder's projects. They are not ClauseGuard-specific. Tone for this project Non-partisan. Cite the bill. Cite the Constitution. Let the evidence make the case. If a finding is contested, say so and present both sides. When writing public-facing reports on sitting government bills, the tone is professional and analytical. Not advocacy. Not opinion. The system's value comes from being trusted by everyone, including the ministry whose bill is being analysed. File structure quick reference What the assistant should do Read docs/methodology/01_taxonomy.md before classifying anything. Read DECISIONS.md before suggesting an architectural change. Cite clause IDs as primary citations. Page numbers only as fallback. Validate every constitutional article cited against the actual text in docs/reference/constitution_1992/articles/. The LLM will invent article numbers if retrieval fails. Validation is a hard rule. docs/methodology/ Risk taxonomy, scoring, labelling guide. The methodological backbone. docs/reference/ 1992 Constitution, indexed by article. docs/bills/ Organised ministry/year/bill. v1 is mocdt/2025/. docs/reports/ Public-facing artefacts. src/clauseguard/ Code. Parsing, retrieval, analysis, evidence, api, eval. DECISIONS.md Architectural decision log. SKILLS.md Capabilities mapped to files and decisions.

When unsure between two risk categories, return both. Multi-label is the default. Preserve clause hierarchy in any output. Clause 12(3)(b) is not the same as clause 12. When generating evidence cards, quote the bill text verbatim. The user verifies against the original PDF. What the assistant should not do Do not invent constitutional articles. Always retrieve before citing. Do not score without evidence. If the retrieval is empty, return "insufficient evidence" rather than a score. Do not generate political analysis or vote predictions. Do not summarise constitutional articles in the output. Quote them. Do not over-engineer the RAG. Use Chroma locally. Use OpenAI or Anthropic APIs for embeddings and LLM calls. Do not build custom vector stores. Do not write code longer than thirty lines per function or one hundred and fifty lines per module. Split when it grows. Do not pre-emptively write tests or modules the founder did not ask for. Build when needed, not in advance. Coding standards Python 3.11. Type hints on every public function. One function does one thing. Functions take inputs as arguments and return outputs. No hidden state. pdfplumber for PDF extraction. spaCy for sentence boundaries. Chroma for vectors. FastAPI for the API. Streamlit for v1 review interface. LLM calls go through a single adapter in src/clauseguard/analysis/llm.py. Easy to swap providers. Every prompt is a versioned file in src/clauseguard/analysis/prompts/. Not hardcoded in Python. Domain quick reference

Bills are hierarchical. Long title, memorandum, parts, clauses, sub-clauses, schedules. The clause splitter must preserve this. The 1992 Constitution is the spine. It is fixed. Articles are stable. Cross-references to other Acts (NCA Act, Data Protection Act, Electronic Communications Act, Cybersecurity Act) are flagged by the parser but not resolved automatically in v1. Risk categories are eleven plus two cross-cutting. Defined in docs/methodology/01_taxonomy.md. Severity and likelihood are ordinal, four bands each. Not numeric scores. When in doubt Ask before building. The founder runs multiple projects in parallel (Govaxis AI, WanderQuest, Voltaic, GreenLens, JPMorgan portfolio projects). Clarifying questions cost less than rework.