

SKILLS

SKILLS.md Maps the technical capabilities demonstrated by ClauseGuard to specific files, decisions, and code paths. Useful for the founder when writing CV bullets or interview talking points. Useful for a recruiter skim-reading the repo to verify each claim. Each capability lists the demonstration, the relevant files, and the relevant decisions in DECISIONS.md. Domain-grounded RAG with a fixed reference corpus Demonstrated by. Two-corpus retrieval where the 1992 Constitution is indexed once as the fixed counter-source and each bill is the variable input. Every clause is tested against the Constitution before any risk is reported. Files. `src/clauseguard/retrieval/embed_constitution.py` `src/clauseguard/retrieval/embed_bill.py` `src/clauseguard/retrieval/retrieve.py` `docs/reference/constitution_1992/` Decisions. D-003. Hierarchical document parsing for legal text Demonstrated by. A clause splitter that preserves the bill's clause, sub-clause, schedule, and cross-reference structure. Output is structured JSON, not flat text. Clause IDs are first-class citizens in the citation system. Files. `src/clauseguard/parsing/clause_splitter.py` `src/clauseguard/parsing/cross_reference.py` `src/clauseguard/parsing/pdf_extractor.py`

Decisions. D-004, D-012. Multi-label classification with an inter-rater agreement protocol Demonstrated by. Each clause is classified across eleven risk categories with no forced single label. Inter-rater agreement is measured against a consulting lawyer before labelling scales beyond thirty clauses. Cohen's kappa threshold is gated at 0.6. Files. `src/clauseguard/analysis/risk_classifier.py` `src/clauseguard/eval/inter_rater.py` `docs/methodology/04_labelling_guide.md` Decisions. D-005, D-010. Transparent ordinal scoring resistant to political pushback Demonstrated by. Severity and likelihood as separate ordinal bands, never collapsed into a single numeric score. No default overall bill score. Stakeholder lenses make the weighting visible by construction. Files. `src/clauseguard/analysis/severity_likelihood.py` `docs/methodology/03_scoring.md` `frontend/lens_selector.py` Decisions. D-006, D-007. Constitutional contradiction checking with hallucination guards Demonstrated by. A contradiction prompt that takes one clause and one constitutional article and outputs an alignment score with rationale. Every cited article is validated against the actual constitutional text before being shown. The model cannot invent article numbers.

Files. `src/clauseguard/analysis/constitutional_check.py` `src/clauseguard/analysis/prompts/constitutional_check.txt` `src/clauseguard/analysis/validators/article_validator.py` Decisions. D-003. Evidence-first system design Demonstrated by. Every score is anchored to a verbatim excerpt from the bill and a verbatim excerpt from the Constitution. The system returns "insufficient evidence" rather than guessing when retrieval is empty. Evidence cards are the unit of output, not scores. Files. `src/clauseguard/evidence/extract_excerpt.py` `src/clauseguard/evidence/format_card.py` `docs/methodology/05_evidence_format.md` Decisions. D-009. Human-in-the-loop ML system architecture Demonstrated by. The v1 workflow is the LLM proposes, the analyst confirms. Edits are captured and feed the labelled dataset. The system is built to grow more autonomous over time as agreement improves, but never starts autonomous in a high-stakes domain. Files. `src/clauseguard/api/review.py` `frontend/review_interface.py` `src/clauseguard/eval/feedback_capture.py` Decisions. D-009.

Extensible repo design for multi-ministry scale Demonstrated by. The `docs/bills/` tree is partitioned ministry first, year next, bill last. Adding a new ministry is a folder operation. The taxonomy, scoring methodology, and Constitution stay shared at the top level. The structure supported by `metadata.yaml` in each bill folder allows querying across the corpus by date, status, or amended Act. Files. `docs/bills/ (structure)` `docs/bills/mocdt/2025/nita_bill/metadata.yaml` `scripts/index_bills.py` Decisions. D-011. Public deliverable in a politically sensitive domain Demonstrated by. The NITA Bill report. Published while the bill is in active parliamentary debate. Non-partisan in tone. Every finding cited to the bill and the Constitution. Distributed through industry, civil society, and parliamentary channels simultaneously. Files. `docs/reports/nita_bill_2025_public.md` The linked LinkedIn post and PDF version Decisions. D-015, D-016. Prompt versioning and provider abstraction Demonstrated by. Every LLM prompt is a versioned file under `src/clauseguard/analysis/prompts/`, not hardcoded in Python. All LLM calls go through a single adapter so providers can be swapped without changing analysis code. Files. `src/clauseguard/analysis/llm.py` `src/clauseguard/analysis/prompts/`

Decisions. Logged in `CLAUDE.md` coding standards. To be promoted to a numbered decision if challenged. How to use this file When writing a CV bullet, find the relevant capability and use the demonstrated language. When preparing for an interview, walk the interviewer through the files and the decisions log entry. When a recruiter asks "show me where you did X", point at the capability section and the linked files. This file is updated whenever a new capability is demonstrated or a code path moves. Last updated on the date of the most recent entry above.